



Hands-On HPC Application Development Using C++ and SYCL

Dounia Khaldi, Hugh Delaney, James Reinders, Michael Wong, Ronan Keryell

HANDLING ERRORS AND DEBUGGING

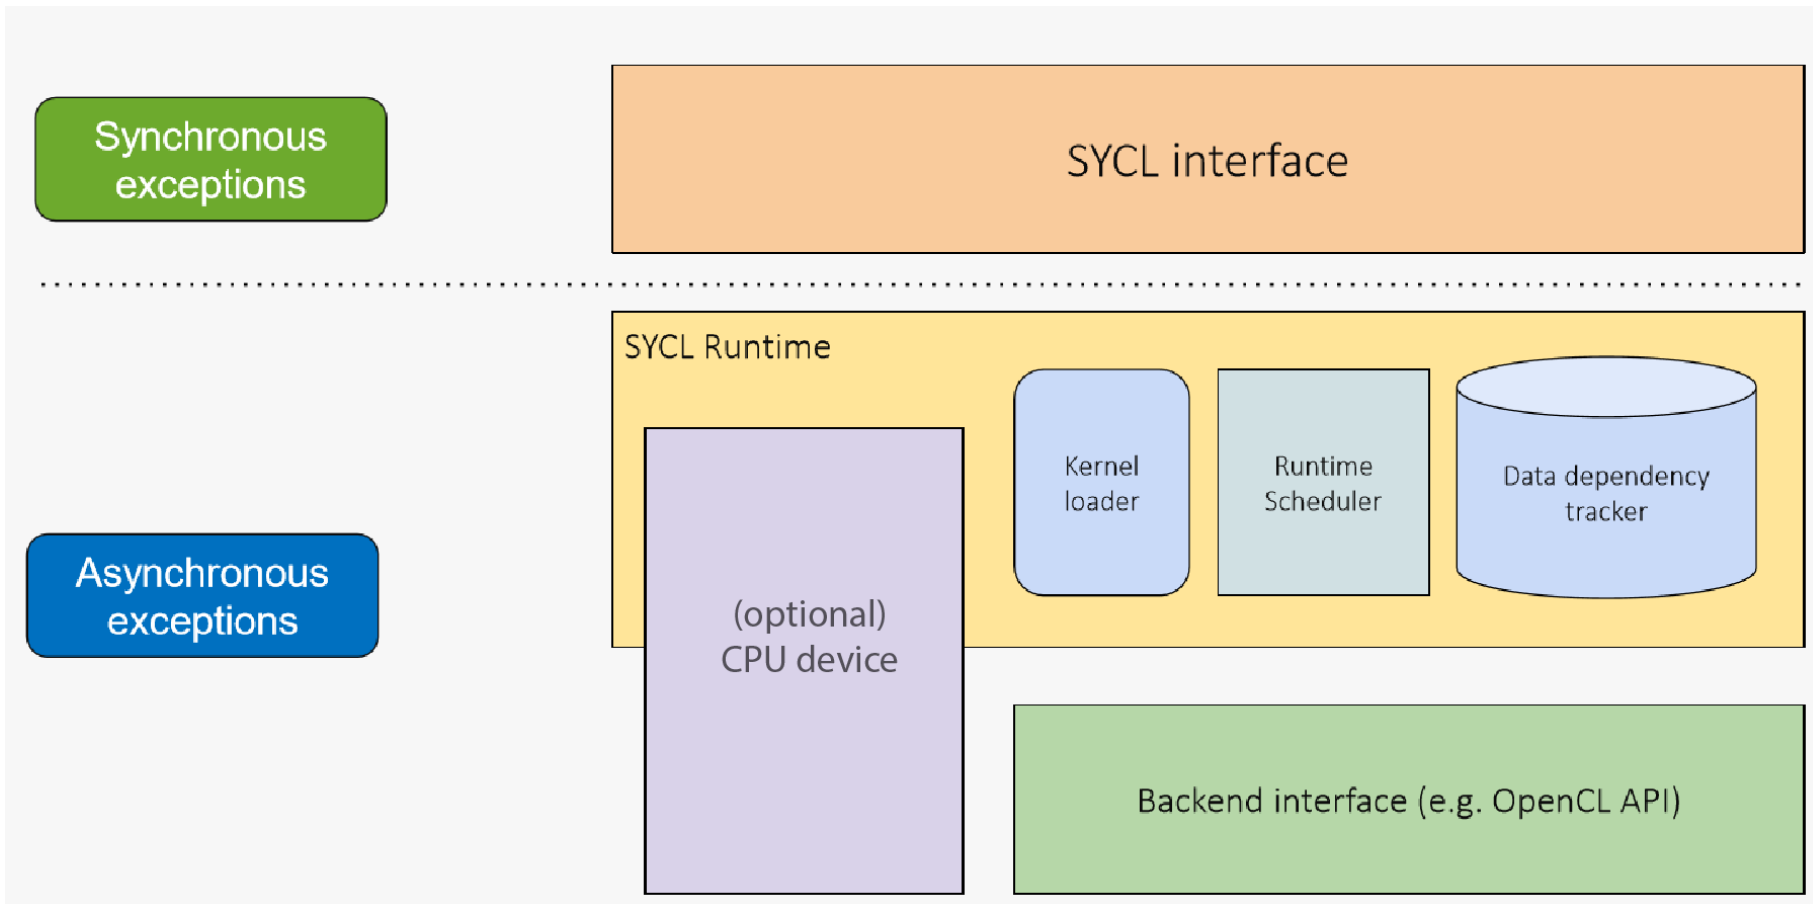
LEARNING OBJECTIVES

- Learn about how SYCL handles errors
- Learn about the difference between synchronous and asynchronous exceptions
- Learn how to handle exceptions and retrieve further information
- Learn about the host device and how to use it

SYCL EXCEPTIONS

- In SYCL errors are handled by throwing exceptions.
- It is crucial that these errors are handled, otherwise your application could fail in unpredictable ways.
- In SYCL there are two kinds of error:
 - Synchronous errors (thrown in user thread) .
 - Asynchronous errors (thrown by the SYCL scheduler).

SYCL EXCEPTIONS



HANDLING ERRORS

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    queue gpuQueue(gpu_selector{});
    buffer bufA{dA};
    buffer bufB{dB};
    buffer bufO{dO};

    gpuQueue.submit([&](handler &cgh) {
        auto inA = accessor{bufA, cgh, read_only};
        auto inB = accessor{bufB, cgh, read_only};
        auto out = accessor{bufO, cgh, write_only};

        cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
            out[i] = inA[i] + inB[i];
        });
    }).wait();
}
```

- If errors are not handled, the application can fail:
 - SYCL 1.2.1 application will fail silently.
 - SYCL 2020 provides a default async handler that will call `std::terminate` when an asynchronous error is thrown.

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try {
        queue gpuQueue(gpu_selector{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

    } catch (...) { /* handle errors */ }
}
```

- Synchronous errors are typically thrown by SYCL API functions.
- In order to handle all SYCL errors you must wrap everything in a try-catch block.

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue gpuQueue(gpu_selector{}, async_handler{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

        gpuQueue.throw_asynchronous();
    } catch (...) { /* handle errors */
    }
}
```



```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
            for (auto e : eL) { std::rethrow_exception(e); }
        });

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        gpuQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        }).wait();

        gpuQueue.throw_asynchronous();
    } catch (...) { /* handle errors */ }
}
```

```
int main() {  
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };  
    try {  
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {  
            for (auto e : eL) { std::rethrow_exception(e); }  
        });  
  
        ...  
  
        gpuQueue.throw_asynchronous();  
    } catch (const std::exception& e) {  
        std::cout << "Exception caught: " << e.what()  
        << std::endl;  
    }  
}
```

- Once rethrown and caught, a SYCL exception can provide information about the error
- The what member function will return a string with more details

```
int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try {
        queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
            for (auto e : eL) { std::rethrow_exception(e); }
        });

        ...

        gpuQueue.throw_asynchronous();
    } catch (const sycl::exception& e) {
        std::cout << "Exception caught: " << e.what();
        std::cout << " With OpenCL error code: "
        << e.get_cl_code() << std::endl;
    }
}
```

- In SYCL 1.2.1, if the exception has an OpenCL error code associated with it this can be retrieved by calling the `get_cl_code` member function
- If there is no OpenCL error code this will return `CL_SUCCESS`
- SYCL 2020 provides the `error_category_for` templated free function that allows checking for the category of the exception depending on the backend used (e.g. `backend::opencl`), and `e.code().value()` will correspond to the backend error code.

```
int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };

    queue gpuQueue(gpu_selector{}, [=](exception_list eL) {
        for (auto e : eL) { std::rethrow_exception(e); }
    });
    context gpuContext = gpuQueue.get_context();

    try {
        ...
        gpuQueue.wait_and_throw();
    } catch (const sycl::exception& e) {
        if (e.has_context()) {
            if (e.get_context() == gpuContext) {
                /* handle error */
            }
        }
    }
}
```

- The `has_context` member function will tell you if there is a SYCL context associated with the error
- If that returns true then the `get_context` member function will return the associated SYCL context object

EXCEPTION TYPES

- In SYCL 1.2.1 there are a number of different exception types that inherit from `std::exception`
 - E.g. `runtime_error`, `kernel_error`
- SYCL 2020 only has a single `sycl::exception` type which provides different error codes
 - E.g. `errc::runtime`, `errc::kernel`

DEBUGGING SYCL KERNEL FUNCTIONS

- Every SYCL 1.2.1 implementation is required to provide a host device
 - This device executes native C++ code but is guaranteed to emulate the SYCL execution and memory model
- This means you can debug a SYCL kernel function by switching to the host device and using a standard C++ debugger
 - For example gdb

- SYCL 2020 only guarantees that a device will always be available, and users can query the `host_debuggable` device aspect to check whether they can use the same functionality as the SYCL 1.2.1 host device

```
class add;

int main() {
    std::vector<float> dA{ 7, 5, 16, 8 }, dB{ 8, 16, 5, 7 }, dO{ 0, 0, 0, 0 };
    try{
        queue hostQueue(aspect_selector<aspect::host_debuggable>(), async_handler{});

        buffer bufA{dA};
        buffer bufB{dB};
        buffer bufO{dO};

        hostQueue.submit([&](handler &cgh) {
            auto inA = accessor{bufA, cgh, read_only};
            auto inB = accessor{bufB, cgh, read_only};
            auto out = accessor{bufO, cgh, write_only};

            cgh.single_task<add>(bufO.get_range(), [=](id<1> i) {
                out[i] = inA[i] + inB[i];
            });
        });
        hostQueue.wait_and_throw();
    } catch (...) { /* handle errors */ }
}
```

QUESTIONS

EXERCISE

Code_Exercises/Exercise_4_Handling_Errors/source

Add error handling to a SYCL application for both synchronous and asynchronous errors.